

Reinforcement Learning Notes from Guillaume Charpiat's Course

Hugo Richard

January 2018

1 Chapter 1: Bandits

1.1 Reinforcement learning framework

The Agent perform an action on the Environment and get a reward that depends on the Environment. In bandits, the environment is static (no environment basically)

Ex: Medicine trial Want to make statistics on medicine A, B, C to know which one is the best. You also want to minimize the number of people dying in the process.

1.2 Strategies

- Greedy: you would try a few times each of them and then you say: "I am done exploring" then you exploit (you take always the best arm in the previous step). Issues: How many times is "a few time" ?
- To be better you need to keep exploring all the time: ϵ -Greedy: At each time-step you flip a coin with probability epsilon you pick randomly uniformly any arm, otherwise (probability $1 - \epsilon$) you take the best arm so far. At initialization you either chose a high or low estimate for each arm (high if you want to force exploration at the beginning)
- Softmax: $Q(a)$ is the average reward for the action a

$$p(a) = \frac{\exp(\frac{Q(a)}{\tau})}{\sum_{a'} \exp(\frac{Q(a')}{\tau})}$$

if τ is close to 0 it tends to be close to the greedy strategy. if τ is close to infinity: you pick every arm with equi-probability.

1.3 Model

k arms to pull. Each time you pull one arm there is a distribution of possible rewards D_k with its value in $[0, 1]$. At time t we chose an arm $I_t \in \{1, \dots, K\}$ and get the reward x_t which follows D_{I_t} (x_t are iid)

We define: $\mu_k = E[x_k]$ and $\mu_* = \max_k(\mu_k)$

The regret for n trials is defined as:

$$R_n = n\mu_* - \sum_{t=1}^n x_t$$

= best-reward-on-average – current-reward

We have

$$\begin{aligned} E[R_n] &= n\mu_* - \sum_t E_{I_t}[\mu_{I_t}] \\ &= E\left[\sum_k T_k(n)(\mu_* - \mu_k)\right] \end{aligned}$$

with $T_k(n)$ the number of time you select the arm k and $\delta_k = \mu_* - \mu_k > 0$.

1.4 UCB: Upper confidence bound

$$B_{t,s}(k) = \widehat{\mu}_{k,s} + \sqrt{\frac{2\log t}{s}}$$

with $s = T_k(t-1)$ the number of time you picked arm k until now and $\widehat{\mu}_{k,s}$ the average reward so far for arm k

At time t , Select:

$$I_t = \operatorname{argmax}_k B_{(t,s)}(k)$$

This is a way of selecting the arm with the highest optimistic average (average + error).

1.5 UCB Bound

The Chernoff-Hoeffding inequality reads:

$$p\left(\frac{\sum_t x_t}{s} - \mu \geq \epsilon\right) \leq \exp(-2s\epsilon^2)$$

So we get

$$P\left(\widehat{\mu}_{k,s} + \sqrt{\frac{2\log t}{s}} \leq \mu_k\right) \leq \exp(-4\log(t)) = 1/t^4$$

Lemma If k is not optimal:

$$E[T_k(n)] \leq 8 \frac{\log(n)}{\delta_k^2} + \frac{\pi^2}{3}$$

We then have

$$\begin{aligned} E[R_n] &= \sum_k \delta_k E[T_k] \\ &\leq \sum_{k \neq k_*} 8 \frac{\log n}{\delta_k} + K \frac{\pi^2}{3} \end{aligned}$$

So $E[R_n]$ is bounded by a linear function of $\log n$.

But δ_k could be very small. This is not a problem because Cauchy-Schwartz yields:

$$\begin{aligned} E[R_n] &= \sum_k \delta_k E[T_k] \\ &\leq \sum_k \sqrt{\delta_k^2 E[T_k]} \sqrt{E[T_k]} \\ &\leq \sqrt{\sum_k \delta_k^2 E[T_k]} \sqrt{\sum_k E[T_k]} \\ &\leq \sqrt{8Kn(\log n + \frac{\pi^2}{3})} \end{aligned}$$

Proof of the lemma Suppose at time t

$$\mu_k - \sqrt{\frac{2\log t}{s}} \leq \widehat{\mu_{k,s}} \leq \mu_k + \sqrt{\frac{2\log t}{s}}$$

Consider k not optimal, k_* the best. It time t , arm k is picked, then

$$\begin{aligned} B_{t, T_k(t-1)}(k) &\geq B_{t, T_{k_*}(t-1)}(k_*) \\ \widehat{\mu_{k,s}} + \sqrt{\frac{2\log t}{s}} &\geq \widehat{\mu_{k_*, s_*}} + \sqrt{\frac{2\log t}{s_*}} \\ \mu_k + \sqrt{\frac{2\log t}{s}} &\geq \mu_* \\ s &\leq 8 \frac{\log t}{\delta_k^2} \end{aligned}$$

$\forall u \in \mathbb{N}$,

$$\begin{aligned} T_k(n) &\leq u + \sum_{t=u+1}^n \mathbb{1}_{I_t=k \wedge T_k(t) \geq u} \\ &\leq u + \sum_{t=u+1}^n \mathbb{1}_{\exists s, s' u < s \leq t \wedge 1 \leq s_* \leq t \wedge B_{t,s}(k) \geq B_{t,s_*}(k_*)} \end{aligned}$$

Finally, consider $u = \frac{8 \log n}{\delta_k^2}$:

$$\begin{aligned} E[T_k(n)] &\leq \frac{8 \log n}{\delta_k^2} + \sum_{t=u+1}^n \left(\sum_{s=u+1}^t t^{-4} + \sum_{s=1}^t t^{-4} \right) \\ &\leq \frac{8 \log n}{\delta_k^2} + \frac{\pi^2}{3} \end{aligned}$$

1.6 Inf Bounds

[Lai et Robbins 1985]

$$\begin{aligned} \limsup_n \frac{E[T_k(n)]}{\log n} &>= \frac{1}{KL(D_k || D^*)} \\ E[T_k(n)] &= \Omega(\log n) \\ E[R_n] &= \Omega(\log n) \end{aligned}$$

[Cesa-Bianchi and Lugos 2006]

$$\inf_{\text{algo}} \sup_{\text{task}} E[R_n] = \Omega(\sqrt{nK})$$

1.7 KL-UCB

$$B_{t,s}(k) = \sup_D \{E[D], ||\widehat{\mu}_{k,s} - E[D]||^2 \leq \frac{f(t)}{T_k(t)}\}$$

D is a model. The term $||\widehat{\mu}_{k,s} - E[D]||^2$ is better replaced with $KL(\widehat{\mu}_{k,s} || D)$.

1.8 Variation on Softmax

$Q_t(k)$ is the total reward of arm k until now.

$$p_t(k) = \exp(Q_t(k))$$

and finally:

$$p(h) = p_t(h)(1 - \gamma) + \gamma 1/K$$

$$\text{If } \gamma = \frac{K \log K}{(e-1)n}$$

$$E[R_n] \leq O(\sqrt{nK \log K})$$

1.9 Combining expert advice

Time series prediction

Each expert gives a distribution of probabilities $D_{k,t}$.

A linear combination of all experts is often better than a single one.

Regret is redefined with respect to the new maximum expected gain.

2 Chapter 2 : Learning dynamics

2.1 Definitions

an Agent interacts with an Environment. The State is given both by the internal configuration of the agent and the environment. Generally speaking a state is a compact description of the world. You can say it is an a-priori where you put all information available. You can also put history in that.

At each step, the agent takes an action, go to a new state and gain a reward. The difference between the agent and the environment is usually defined by what you can control and what you cannot.

When you are in a particular state you chose an action based on a policy. We try to find the best policy to get maximum reward.

Example 1 Let us consider a Robot with an arm. The position of the arm represents the state of the robot.

An agent may not know the full state of the world. It may have only a partial view.

Example 2 Let us consider a Robot moving. The agent is the controller, the part making decision. It is not the hardware.

Example 3 Let us consider a chess game. If you specify sub-goals, for example gaining points when taking adverse pieces, you might end up taking the queen and losing the game: subgoals are dangerous.

The goal is to maximize rewards on the long run.

Let us consider a finite horizon T (finite number of turn in a game). At time t you want to maximize:

$$E_{e,\Pi} \left[\sum_{i=t+1}^T r_i \right]$$

where Π is the policy, a_t , s_t , r_t are respectively action, state and reward at time t and e is the environment.

In practice there might not be a specified finite number of turn but there usually exists terminal state (states that end the game).

You might want to put less weight on rewards you will gain in a very far future.

$$E_{e,\Pi} \left[\sum_{k=0}^{\infty} \gamma^k r_{t+1+k} \right]$$

where γ , the discount rate is such that $0 \leq \gamma \leq 1$.

We define the return R by:

$$R = \sum_{k=0}^{\infty} \gamma^k r_{t+1+k}$$

2.1.1 Markov Decision Process (MDP)

To choose an action you might want to evaluate the quantity:

$$p = p(s_{t+1}, r_{t+1} | s_t, a_t, r_t, s_{t-1}, a_{t-1}, r_{t-1}, \dots)$$

We make the hypothesis of a Markovian environment (a way to do that is to include history but remember we want short states...) so that we have $p = p(s_{t+1}, r_{t+1} | s_t, a_t)$

In a Markov Decision Process, everything is described by only two quantities:

- $p(s_{t+1} | s_t, a_t)$: this describes the dynamics (how we are moving in the space of state)
- $E_e[r_{t+1} | s_t, a_t, s_{t+1}]$: the average reward (full distribution not required)

This can be described as a graph (like a Markov chain).

2.1.2 Value functions: V and Q

Let us suppose that an environment follows a Markov Decision Process and that we know all the state and reward graph.

Policy Π is a function such that $\Pi(state, action)$ is the probability of taking that action in that state. Of course $\sum_a \Pi(state, action) = 1$.

Let us assume you are in the state s_t you chose action a_t using policy distribution and a new state and reward is sampled from the environment.

A State value function V is a function such that $V(state)$ is the expected return when you are in that state.

$$V^\Pi(s) = E_{\Pi, e}[R_t | s_t = s] = E_{e, \Pi}[\sum_{k=0}^{\infty} \gamma^k r_{t+1+k} | s_t = s]$$

with s_T the terminal state. We have $V^\Pi(s_T) = 0$.

An State-action value function Q is a function such that $Q(state, action)$ is the expected return when you are in that state and you take that action.

$$Q^\Pi(s, a) = E_{\Pi, e}[R_t | s_t, a_t]$$

Consider any policy Π and any state s . At each time step we chose an action a_t which depends on the policy Π and the current state s_t . Depending on the action a_t the current state s_t and the environment e , we get a new state s_{t+1} and a reward r_{t+1} . We denote $E_{e, \Pi}$ the expectation taken on all possible actions,

states and rewards.

$$\begin{aligned}
V^\Pi(s) &= E_{\Pi,e}[R_t | s_t = s] \\
&= E_{e,\Pi}[\sum_{k=0}^{\infty} \gamma^k r_{t+1+k} | s_t = s] \\
&= E_{e,\Pi}[r_{t+1} + \sum_{k=0}^{\infty} \gamma^k r_{t+2+k} | s_t = s] \\
&= E_{a_t \sim \Pi, s_{t+1} \sim e}[r_{t+1} + \gamma V^\Pi(s_{t+1})]
\end{aligned}$$

2.1.3 Optimal Policies

Let us suppose we have a markov decision process and that there is a partial order over policies. We say $\Pi \geq \Pi' \iff \forall s, V^\Pi(s) \geq V^{\Pi'}(s)$.

Properties

- Existence of Optimal policy:

$$\exists \Pi^*, \forall \Pi, \Pi^* \geq \Pi$$

- All possible Π^* have the same value functions:

$$\forall s, Q^*(s, a) = \max_{\Pi} Q^\Pi(s, a)$$

Bellman optimality equation

$$V^*(s) = \max_a \sum_{s'} p(s'|s, a) (R_{s,s'}^a + \gamma V^*(s'))$$

$R_{s,s'}^a$ is the average reward when taking action a in state s and arriving in s' and $p(s'|s, a)$ depends only on the environment.

Similarly:

$$Q^*(s, a) = \sum_{s'} p(s'|s, a) (R_{s,s'}^a + \gamma \max_{a'} Q^*(s', a'))$$

Optimizing the policy When you have a policy Π you want to find the value function V^Π (this is called policy evaluation)

First step: Estimating the value function An iterative approach works as follow. Let start with any value function V_k .

$$\begin{aligned}
V_{k+1}(s) &= E_{\Pi,e}[r_{t+1} + \gamma V_k(s_{t+1}) | s_t = s] \\
&= \sum_a \pi(s, a) \sum_{s'} p(s'|s, a) (R_{s,s'}^a + \gamma V_k(s'))
\end{aligned}$$

I am in state s_t . I need to take an action a_t . I arrive in a new state (s_{t+1}, r_{t+1}) with some probabilities. My estimation of r_{t+1} is exact but the estimation of $V_k(s_{t+1})$ is the one I assumed before (at iteration k). This converges to the real value function since $\gamma < 1$. Making use of $V_k(s_{t+1})$ to re-estimate $V_{k+1}(s_t)$ is named bootstrapping.

The proof of convergence is given using

$$\begin{aligned} |V_{k+1} - V^\Pi(s)| &= \left| \sum_a \pi(s, a) \sum_{s'} p(s'|s, a) [R_{s,s'}^a + \gamma V_k(s') - R_{s,s'}^a - \gamma V^\Pi(s')] \right| \\ &\leq \sum_a \pi(s, a) \sum_{s'} p(s'|s, a) \gamma |V_k(s') - V^\Pi(s')| \\ &\leq \sum_a \pi(s, a) \sum_{s'} p(s'|s, a) \gamma \|V_k - V^\Pi\|_\infty \end{aligned}$$

$$a \leq \gamma \|V_k - V^\Pi\|_\infty$$

We define the infinite norm by $\|V_{k+1} - V^\Pi\|_\infty = \max_s |V_{k+1} - V^\Pi(s)|$ and we find

$$\|V_{k+1} - V^\Pi\|_\infty \leq \gamma \|V_k - V^\Pi\|_\infty$$

There exist an inplace version where update for s are done in series rather than in parallel.

Policy improvement We have some policy Π we describe how good it is with V^Π and Q^Π . Define greedy policy Π' which consist in taking the best action according to the previous value function

$$\Pi'(s) = \operatorname{argmax}_a Q^\Pi(s, a)$$

Thus $\Pi'(s) \geq \Pi(s)$

If no improvement the policy is optimal already. Doing policy evaluation and policy improvement repetitively is called policy iteration.

Let us introduce T the Bellman operator such that $T(\Pi) = \Pi'$ is the result of policy improvement algorithm on Π .

The optimal policy Π^* satisfies $T\Pi^* = \Pi^*$ which we can write $B\Pi^* = 0$ with $B = T - Id$.

The optimization we do is equivalent to Newton Method.

$$\Pi_{k+1} = \Pi_k + (dB/d^2\Pi)^{-1} dB/d\Pi(\Pi_k)$$

A method called truncated policy evaluation works as follow. You don't wait for convergence, you just apply one step of policy evaluation before doing policy improvement. Still same guaranties of convergence. The proof is the same.

This can be generalized to any number of steps, the same guaranties always apply. With this dynamic programming algorithm, finding optimal policy is polynomial in the number of actions m and states s . Without this trick (naive evaluation of all policies) we get exponential complexity $O(m^n)$.

2.2 Monte Carlo (MC) methods

The value function $V^\Pi(s)$ is the average return in state s .

When running simulation, we are playing episodes = runs of the game. Averaging the empirical value function for a state during many episodes (one episode = one sample) might be a good representation of the value function.

Whether you consider only the first visit or every visit to the state, your estimate is always going to converge to $V^\Pi(s)$ with precision $\frac{1}{\sqrt{n_{samples}}}$.

We have a fixed policy Π you sample using Monte Carlo Q and you update your policy using the greedy policy improvement presented in last section.

2.2.1 On-Policy MC-control

We sample according to the policy we want to use. An additional hypothesis is that you have ϵ -soft policy with $\epsilon > 0$. This means $\forall s, a \Pi(s, a) > \epsilon$ (ϵ -greedy is one example of such policy).

2.2.2 Out-Policy evaluation

We can evaluate Π based on experiments with Π' . It requires that $\Pi(s, a) > 0 \implies \Pi'(s, a) > 0$. Consider the first visit to a state s . Let us denote p_i the likelihood of future state, reward and action for policy Π and let us denote p'_i the likelihood of future state, reward and action for policy Π' .

$$p_i(s) = \prod_{k=t}^{T-1} \pi(s_k, a_k) p(s_{t+1} | a_k, s_k)$$

$$p'_i(s) = \prod_{k=t}^{T-1} \pi'(s_k, a_k) p(s_{t+1} | a_k, s_k)$$

Instead of averaging like $V(s) = 1/n \sum_i R_i(s)$, we consider

$$V(s) = \frac{1}{Z} \sum_i \frac{p_i(s)}{p'_i(s)} R_i(s)$$

$$= \frac{1}{Z} \sum_i \frac{\pi(s_k, a_k)}{\pi'(s_k, a_k)} R_i(s)$$

The behaviour policy is different than the estimation policy. The behaviour policy is used for sampling (running experiments). The estimation policy is the policy you really want to use in the end (the one you want to improve).

```

forall s, a
  Q(s, a) ← random()
  N(s, a) ← 0
  D(s, a) ← 0
  Π ← random()
repeat

```

generate an episode $s_0, a_0, r_0, s_1, a_1, r_1, \dots$ with any soft policy Π'
 take the tail ($t > \tau$) so that all $a_{t>\tau}$ follows $\Pi(s_t)$
 $\forall (s, a)$ in that tail
 $w \leftarrow \prod_{k=1}^{T-1} \frac{1}{\Pi'(s_k, a_k)}$
 $N(s, a)+ = wR_t$
 $D(s, a)+ = w$
 $Q(s, a) = \frac{N}{D}(s, a)$
 $\forall s, \Pi(s) = \operatorname{argmax}_a Q(s, a)$
until enough iteration are made

On the good size of MC methods, there is no need to estimate environment bootstrap or probabilities. However, many exploration might be needed.

2.2.3 Temporal difference

Immediate MC on reward (not returns).

$$V(s_t)+ = \alpha(r_{t+1} + \gamma(V(s_{t+1}) - V(s_t)))$$

with $\alpha \in [0, 1]$

This is bootstrapping (using $V(s_{t+1})$)

Proved to converge to V^Π if α is small enough or $\alpha_t \rightarrow 0$ and $\sum \alpha_t = \infty$

You can also do this in batch mode: You run a number of episodes and then update (faster convergence).

2.2.4 On-policy TD control: SARSA

We have a policy Π we want to estimate the value $Q(s, a)$ (same as before replacing V with Q)

$$Q(s_t, a_t)+ = \alpha[r_{t+1} + \gamma(Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t))]$$

Update for each 5-uple $(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})$ hence the name SARSA.

Π derived from Q needs to be soft e. g. ϵ -greedy.

2.2.5 Off-policy TD control: Q-learning

Breakthrough at that time because it is working better than other approaches.

- Exploitation

$$Q_t(s_t, a_t)+ = \alpha(r_{t+1} + \gamma \max_a (Q(s_{t+1}, a)) - Q(s_t, a_t))$$

- Exploration: Choose an action a using a policy derived from Q e.g ϵ -greedy (for softness).